

Artificial Atoms: Modeling the Inductively-Shunted Cooper Pair Box

John O'Connor

The Cooper pair box (CPB) is a well-studied quantum circuit that can function as an artificial atom. The inductively-shunted Cooper pair box (CPBL) is much less well-studied. Adding an inductive shunt reduces the sensitivity of the device to ambient charge noise around the circuit, but it also produces a Hamiltonian with completely different characteristics. The first examples of these circuits are being studied in Michel Devoret's lab by Vladimir Manucharyan. While computer models of the energy spectrum for these circuits exist, faster algorithms would be helpful in interpreting the experimental data more efficiently. We develop a computer model of the CPBL that is fast enough to run curve fitting on experimental data, and demonstrate the results of this model.

I. INTRODUCTION

Quantum computing offers a number of enticing theoretical speedups over classical computers. Best known is Peter Shor's factoring algorithm[7], which can factor large composite numbers in polynomial time. This has the potential to break cryptographic schemes like RSA that rely on the intractability of the factoring problem. Another less dramatic but more widely applicable speedup is Grover's algorithm for searching an arbitrary space of n elements in time $\mathcal{O}(n)$, a quadratic improvement over the classical case[5]. These and other quantum algorithms have motivated significant research in the field of quantum computing.

The greatest obstacle to constructing a quantum computer has been creating the quantum bits (qubits) that both store information reliably and allow that information to be manipulated. Any quantum system tends to decohere due to interactions with its environment, and a qubit has to meet two contradictory requirements. First, it must be sufficiently decoupled from its environment to hold its state long enough to perform meaningful computation. At the same time, though, it has to be coupled to its environment enough to be manipulable and eventually measurable with the computation is finished. This is difficult enough that, at present, one of the largest quantum computations ever performed was the factorization of the number fifteen.[8]

One of the fields in which research is proceeding on this front is circuit quantum electrodynamics[1][10]. Quantum circuits have several potential advantages over other forms of qubits. They are permanently stationary, which greatly simplifies research compared to various atomic methods, and they can be manufactured using standard circuit lithography techniques. More importantly, while atomic systems have fixed energy spectra, researchers can tune the parameters of their quantum circuits to produce artificial atoms with a range of different spectra.

Computer models of these circuits can aid experimental research in several ways. Theoretical models of the device spectrum highlight extraneous or unexpected features in the data, and help to identify those features. Computer models can explore parameter regimes not yet probed by experiment, and they can do it in seconds rather than days or weeks. Also, the manufacturing pro-




FIG. 1: A diagram of the Cooper Pair Box. Note that E_C and E_J are both parameters of the Josephson junction.




FIG. 2: A sketch of a Josephson junction

cesses for our circuits typically involve significant uncertainties. The insulating layer in a Josephson junction, for example, is formed by allowing an exposed layer of aluminum to oxidize. The tunneling strength of that layer is exponentially sensitive to its thickness, and the layer continues to grow even after the circuit is complete. For this and similar reasons, the design parameters of our circuits are only known approximately until we measure them against theoretical models.

II. MODELING THE COOPER PAIR BOX

One well-studied example of a quantum circuit is the Cooper Pair Box (CPB)[4][9]. See Fig. 1. We used the basic CPB as an initial test of our modeling techniques, and we use it here as an introduction to the inductively-shunted Cooper pair box (CPBL). The CPB is a superconducting circuit with three components: a voltage source, a gate capacitor, and a Josephson junction.

The Josephson junction (Fig. 2) is a circuit element consisting of two superconducting leads joined across a thin insulating boundary. In our devices, Josephson junc-

tions are constructed by depositing a layer of aluminum, allowing a thin oxide layer to form on the surface of that layer, and finally depositing a second layer of aluminum on top. The oxide acts as an insulator, and at sufficiently low temperatures the aluminum leads become superconducting. A Josephson junction acts like a capacitor across its insulating layer, but it also allows Cooper pairs of electrons to tunnel through. We can treat the Josephson junction as a capacitor and a tunneling junction in parallel, and we denote the energy scales corresponding to these components as E_C and E_J , respectively.

The Josephson tunnel junction, together with the gate capacitor, creates a superconducting island in the CPB, indicated by the dotted line in Fig. 1. This lets us treat the CPB in its “charge basis,” where each state denotes a number of Cooper pairs added to or removed from the island. In the charge basis, the Hamiltonian of the CPB is

$$H_{\text{CPB}} = \sum_{n=-\infty}^{+\infty} \left(\frac{-E_J}{2} (|n\rangle \langle n-1| + |n-1\rangle \langle n|) + 4E_C(n - N_g)^2 |n\rangle \langle n| \right), \quad (1)$$

where n represents the number of tunneled Cooper pairs and N_g is the offset charge across the capacitor, which is set by the gate voltage according to $N_g = \frac{C_g V_g}{2e}$.

In order to calculate the energy spectrum of the CPB as a function of the offset charge, we will need to express H_{CPB} as a matrix in the charge basis. Note that H_{CPB} will be an infinite matrix not just downward and rightward but, because n and m can take all integer values, leftward and upward as well. Each matrix element $H_{\text{CPB}nm}$ is equal to the inner product $\langle n | H_{\text{CPB}} | m \rangle$, so specifying those inner products specifies the matrix itself. Eqn. (1) gives us:

$$\langle n | H_{\text{CPB}} | m \rangle = \begin{cases} 4E_C(n - N_g)^2 & \text{if } n = m \\ \frac{-E_J}{2} & \text{if } |n - m| = 1 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

Truncating this tridiagonal matrix allows us to approximate its eigenvalues numerically. Up to an overall scale, the eigenvalues depend only on the ratio between E_J and E_C , and we can plot the energy spectra relative to N_g for several values of this ratio in Figs. 3, 4, and 5. Units on the y-axis are in gigahertz. We can see that for values where E_C is dominant (Fig. 3) the spectrum looks like a series of parabolas, which results from the quadratic term on the diagonal of H_{CPB} . For larger values the spectral lines smooth out and move apart, approaching a constant spacing.




FIG. 3: Energy spectrum vs. N_g for the CPB with $\frac{E_J}{E_C} = 1$




FIG. 4: Energy spectrum vs. N_g for the CPB with $\frac{E_J}{E_C} = 10$

III. CHALLENGES WITH THE CPB

As with all quantum circuits, one of the primary goals is to obtain a device with long decoherence times. In the case of the CPB, one of the primary causes of decoherence is the charge noise—slow, randomly oscillating charges in the environment that introduce noise into the offset charge, N_g .

Several strategies exist for mitigating the effects of charge noise in the CPB. One approach is to operate the device near local extrema—the so-called “sweet spots”—of the energy spectrum[9]. We can see in Fig. 3 that at certain values of N_g all of the spectral lines of the CPB reach local maxima or minima. When the CPB is operated at these points, the first order effects of charge noise drop out.

Another approach is to operate the CPB in the “transmon” region where $EJ \gg EC$ [6]. We can see in




FIG. 5: Energy spectrum vs. N_g for the CPB with $\frac{E_J}{E_c} = 30$. Note that the deviations from periodicity on the left and right ends of each line are artifacts of the approximation process, not of the theory itself. Extending the matrix cutoff pushes these deviations farther out in N_g .




FIG. 6: A diagram of the inductively shunted Cooper pair box (CPBL). The circuit now has three parameters, E_C and E_J associated with the Josephson junction, and E_L associated with the inductor.

Fig. 5 that the spectral lines become very flat, reducing the effects of charge noise even further. One of the disadvantages of working in the transmon region is that, as we mentioned earlier, the spacing between spectral lines becomes nearly uniform. Because a qubit generally needs to be restricted to two energy states, this lack of anharmonicity can be a substantial problem.

A third option—the option that we explore here—is to add an inductive shunt to the CPB (Fig. 6). The energy spectrum of the CPBL is entirely distinct from that of the CPB and worthy of study in its own right, but it is also much less sensitive to charge noise. This results from the fact that at low frequencies—that is, at the frequencies of substantial charge noise—the inductor acts like a closed circuit and shunts those signals, while at high frequencies—like those of our input voltages—the inductor approximates an open circuit, and the CPBL behaves similarly to the CPB.

IV. MODELING THE INDUCTIVELY SHUNTED COOPER PAIR BOX

Adding the inductor requires that we make several changes to our equations. First, since no superconducting island exists in the CPBL, the charge basis is no longer meaningful. We will instead use another basis from the CPB called the phase basis. The phase term arises from the Josephson equations (see [3]). Expressed in the phase basis, the Hamiltonian of the CPB takes the form

$$H_{\text{CPB}} = -4E_C \frac{d^2}{d\varphi^2} - E_J \cos \varphi. \quad (3)$$

Another consequence of the missing island is that no offset charge can exist on the CPBL, and so we will instead plot the energy spectrum with respect to ϕ , the magnetic flux through our inductor. The ϕ parameter comes into the picture when we add the term for our inductor to H_{CPB} , producing the Hamiltonian

$$H = -4E_C \frac{d^2}{d\varphi^2} - E_J \cos \varphi + \frac{1}{2} E_L \left(\varphi - \frac{2\pi\phi}{\phi_0} \right)^2. \quad (4)$$

This equation also introduces a third design parameter for the CPBL, the inductive energy, E_L . In order to analyze this equation, we will immediately adopt a change of variables, $\varphi \rightarrow (\varphi + \frac{2\pi\phi}{\phi_0})$.

$$H = -4E_C \frac{d^2}{d\varphi^2} - E_J \cos \left(\varphi + \frac{2\pi\phi}{\phi_0} \right) + \frac{1}{2} E_L \varphi^2. \quad (5)$$

After the change we can see that, apart from the cosine term in the middle, H takes the form of a simple harmonic oscillator (SHO). It will therefore be convenient to express H in the harmonic oscillator basis over φ . Note that this is again a discrete basis, though unlike the CPB, its matrix form extends infinitely only in the two conventional directions. To facilitate calculations in this basis, we separate H into two pieces:

$$H = H_0 + V \quad (6)$$

$$H_0 = -4E_C \frac{d^2}{d\varphi^2} + \frac{1}{2} E_L \varphi^2 \quad (7)$$

$$V = -E_J \cos \left(\varphi + \frac{2\pi\phi}{\phi_0} \right). \quad (8)$$

Simply substituting the coefficients in H_0 into the SHO equations, with $|n\rangle$ now referring to SHO basis states over φ , gives us that

$$\langle n | H_0 | m \rangle = \begin{cases} \sqrt{8E_L E_C} (n + \frac{1}{2}) & \text{if } n = m \\ 0 & \text{otherwise.} \end{cases} \quad (9)$$

The inner product with V , on the other hand, requires a




FIG. 7: Energy spectrum (difference from the ground state) for the CPBL plotted against magnetic flux. $E_C = 2.5$, $E_J = 8.8$, and $E_L = 0.5$. The y-axis and parameter values are given in gigahertz, and the x-axis is given in multiples of the base flux, ϕ_0 , over one full period.

fairly complicated integral. The result is

$$\langle n | V | m \rangle = \begin{cases} -E_J (-2)^{-m'} \cos\left(\frac{2\pi\phi}{\phi_0}\right) \varphi_0^{2m'} \\ \quad \times \sqrt{\frac{n'!}{(n'+2m')!}} \exp\left(\frac{-\varphi_0^2}{4}\right) \\ \quad \times L_{n'}^{2m'}\left(\frac{\varphi_0^2}{2}\right) & \text{if } n \equiv m \pmod{2} \\ -E_J (-2)^{-m'} \frac{1}{\sqrt{2}} \sin\left(\frac{2\pi\phi}{\phi_0}\right) \varphi_0^{2m'+1} \\ \quad \times \sqrt{\frac{n'!}{(n'+2m'+1)!}} \exp\left(\frac{-\varphi_0^2}{4}\right) \\ \quad \times L_{n'}^{2m'+1}\left(\frac{\varphi_0^2}{2}\right) & \text{if } n \not\equiv m \pmod{2}, \end{cases} \quad (10)$$

where

$$n' = \min(n, m) \quad (11)$$

$$m' = \left\lfloor \frac{|n - m|}{2} \right\rfloor \quad (12)$$

$$\varphi_0 = \sqrt[4]{\frac{8E_C}{E_L}}, \quad (13)$$

and L_n^α is the generalized Laguerre polynomial.

Eqn. (10) is a ghastly expression but straightforward to compute, so it and Eqn. (9) specify H entirely. As we did with the CPB, we can now truncate this Hamiltonian and approximate its eigenvalues numerically. Again, the y-axis and parameter values are given in gigahertz. Fig. 7 shows the spectrum calculated as a difference from the ground state and plotted with respect to magnetic flux, using a set of parameter values close to those that our lab has tested experimentally.

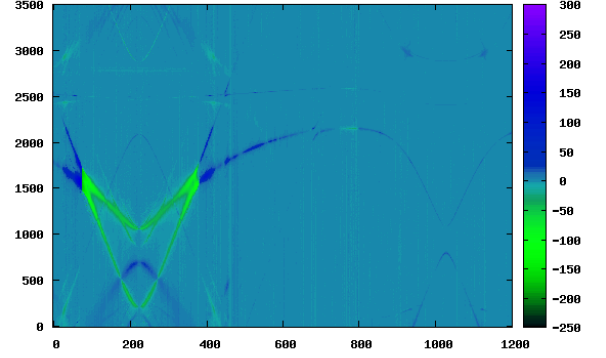


FIG. 8: Raw data. For scale, the near intersection on the bottom right of the image corresponds to the near intersection in the center of Fig. 7.




FIG. 9: Curves derived from the data in Fig. 8

V. EXPERIMENTAL DATA

Vladimir Manucharyan has been measuring the energy spectra of the first CPBL device ever produced, and one of the goals of this project is to fit our theoretical spectra against his data. Energy curves were originally drawn from those data by hand, but a simple threshold filter can automate that process. Raw data (Figs. ?? and 8) can be filtered into sets of points suitable for fitting (Figs. ?? and 9).

There are many ways to define a “badness of fit” for theoretical curves against these data. In our case, the simplest and most effective method is, for each point in the data set, to find the vertically nearest curve in the theoretical set, take the difference between it and the data point, and add together the squares of these differences. That is, for a given fit the badness, B , is defined




FIG. 10: An initial guess for parameter values plotted against experimental data from a portion of the set depicted in Fig. 8. This initial guess is clearly very poor, even though none of the parameters is off by more than 20%.




FIG. 11: The same data set from Fig. 10 after our optimization routines have adjusted the parameter values from the initial guess.

as

$$B = \sum_{p \in \text{data}} (t(p) - p)^2, \quad (14)$$

where $t(p)$ is the closest point in the theoretical data set with the same flux. We can then use generic multi-parameter minimization routines to find the best fitting values of E_C , E_J , and E_L [2][11]. The results of this approach are displayed in Figs. 10 and 11.

The only additional input to the fit depicted in Fig. 11 is the initial parameter guess and the identification of the zero-flux point on the x-axis. Note that the lines jutting out to the left and right of the bottom peak in that figure represent a two-photon transition, and the fit

could be improved by including those transitions. The line appearing clearly between the upper two spectral lines in that figure is not one of these transitions but rather a resonance from our experimental setup outside the CPBL itself, and it causes the nearby lines to deviate slightly from the theory. Cleaning out these extraneous points would improve the fit as well, though it would be difficult to automate.

VI. SPEEDING IT UP

Simulation is one of the few areas left in programming where computational speed remains a limiting factor. Mathematica models of the CPBL built by Jens Koch already existed when we began this project, and we used them to verify our results. They were fast enough to generate the plots shown here but not fast enough to run that computation over and over again to perform a fit. We used several different optimizations to achieve the speeds necessary to perform these fits.

The process of fitting a data set requires repeatedly calculating energy spectra for different values of E_C , E_J , and E_L . Calculating those spectra means repeatedly solving for the eigenenergies of H across a range of values of ϕ values. Naturally, lots of work is duplicated when these calculations are all performed in isolation. As it turns out, the most expensive operation in our case is calculating the coefficients of the generalized Laguerre polynomials to construct the H matrix. (One might guess that solving for eigenvalues would be harder, but since we only need a handful of eigenenergies we can keep our matrices fairly small.) Since Laguerre polynomials obey a recursive relation, some time can be saved by solving for all of them at once, rather than separately for each element of H . However, a far more important realization is that the coefficients of L_n^α are independent of both ϕ and the experimental parameters, so the matrix of polynomials only needs to be calculated once for an entire fitting procedure.

A similar speedup results from the fact that in all of Eqns. 9 and 10, only two terms in Eqn. (10) depend on the value of ϕ . Thus the majority of work needed to assemble H can be done once for each trial of the parameters E_C , E_J , and E_L , rather than for every value of ϕ plotted within that trial. That leaves only one trigonometric term (either $\cos(2\pi\phi/\phi_0)$ or $\sin(2\pi\phi/\phi_0)$) to be calculated and the diagonalization operation to be performed for each value of ϕ . These optimizations altogether speed up a fitting operation by almost twenty times.

We achieved a second major speedup by taking advantage of the capabilities of modern optimization algorithms[2][11]. Many of these algorithms calculate local partial derivatives in order to improve their search for minima. Without assistance, these routines can approximate local derivatives by evaluating several points in close proximity, but it is also possible to supply these

routines with local derivative information we calculate ourselves. Rather than having our routines evaluate the badness of a fit numerous times in a small region around each set of trial parameters, we can calculate the partial derivatives of the badness of fit in parallel with the fit itself. As one might expect, calculating the derivatives this way avoids lots of duplicated effort.

Recalling our definition of “badness of fit” in Eqn. (14), the partial derivative for a given parameter E_* is

$$\frac{\partial B}{\partial E_*} = \sum_{p \in \text{data}} 2(t(p) - p) \frac{\partial t}{\partial E_*}. \quad (15)$$

The values of $t(p)$ are just differences among eigenvalues of H , so we need to calculate the derivatives of those eigenvalues. Just how to do that is not immediately obvious, but beginning with Schrödinger’s equation, we can develop a very convenient intermediate result.

$$H |\psi\rangle = E |\psi\rangle \quad (16)$$

Take the partial derivative of both sides.

$$\frac{\partial}{\partial E_*} (H |\psi\rangle) = \frac{\partial}{\partial E_*} (E |\psi\rangle) \quad (17)$$

Expand via the product rule.

$$\frac{\partial H}{\partial E_*} |\psi\rangle + H \frac{\partial}{\partial E_*} |\psi\rangle = \frac{\partial E}{\partial E_*} |\psi\rangle + E \frac{\partial}{\partial E_*} |\psi\rangle \quad (18)$$

Multiply both sides by $\langle\psi|$.

$$\begin{aligned} \langle\psi| \frac{\partial H}{\partial E_*} |\psi\rangle + \langle\psi| H \frac{\partial}{\partial E_*} |\psi\rangle &= \\ \langle\psi| \frac{\partial E}{\partial E_*} |\psi\rangle + \langle\psi| E \frac{\partial}{\partial E_*} |\psi\rangle \end{aligned} \quad (19)$$

The H in the second term comes out as a factor of E , and the third term simplifies.

$$\langle\psi| \frac{\partial H}{\partial E_*} |\psi\rangle + E \langle\psi| \frac{\partial}{\partial E_*} |\psi\rangle = \frac{\partial E}{\partial E_*} + E \langle\psi| \frac{\partial}{\partial E_*} |\psi\rangle \quad (20)$$

Finally, the two common terms cancel.

$$\frac{\partial E}{\partial E_*} = \langle\psi| \frac{\partial H}{\partial E_*} |\psi\rangle. \quad (21)$$

Eqn. (21) leaves us with a very convenient formula for calculating the partial derivatives of the eigenenergies of H .

Now, going back to Eqn. (5), we see that

$$\frac{\partial H}{\partial E_C} = -4 \frac{\partial}{\partial \varphi^2} \quad (22)$$

$$\frac{\partial H}{\partial E_J} = \cos(\varphi + \frac{2\pi\phi}{\phi_0}) \quad (23)$$

$$\frac{\partial H}{\partial E_L} = \frac{1}{2} \varphi^2. \quad (24)$$

Our eigenvalue routines can give us the eigenvectors $|\psi\rangle$ associated with each eigenvalue E with little additional effort. In the case of $\partial H / \partial E_J$, we have already calculated the matrix form of Eqn. (23), namely the matrix V without the extra factor of $-E_J$, and taking the inner product of that matrix with $|\psi\rangle$ gives us $\frac{\partial E}{\partial E_J}$.

In order to deal with Eqns. 22 and 24, we return to the machinery of the SHO and recognize that

$$-4 \frac{\partial}{\partial \varphi^2} = -\sqrt{\frac{E_L}{2E_C}} (a - a^\dagger)^2 \quad (25)$$

and

$$\frac{1}{2} \varphi^2 = \sqrt{\frac{E_C}{2E_L}} (a + a^\dagger)^2. \quad (26)$$

We could express these operators in matrix form, but it is simpler and faster to just apply the raising and lowering operators explicitly to each eigenvector and dot the result with the original eigenvector.

Finally, we just sum up the partial derivatives using Eqn. (15) just as we sum up the result with Eqn. (14). Calculating the derivatives ourselves, rather than approximating them with extra function calls, nets us another speedup of nearly a factor of four.

VII. CONCLUSION

All the code for our project was written in Python, making extensive use of the SciPy libraries. While Python is an excellent language for fast prototyping, Python programs have a reputation for running slowly, and at the beginning of our project it was unclear whether we could realize substantial gains in efficiency without eventually moving to a compiled language like C. The fact that we achieved efficiency gains of two orders of magnitude using Python alone demonstrates that there was significant room for improvement both in the algorithmics and in the mathematics of the CPBL model. At this point, our code can fit a complete data set in one to two minutes, and thinning out the data can bring that time down to about fifteen seconds without significantly affecting the accuracy of the fit. We hope that our code will be useful to Vladimir Manucharyan and others as they continue to explore the CPBL, and that the techniques we developed will prove useful for implementing future models of the CPBL and other systems.

VIII. ACKNOWLEDGEMENTS

Thanks to Steven Girvin, Luigi Frunzio, and Lev Bishop for discussions about this project, to Vladimir Manucharyan, Rob Schoelkopf, and Michel Devoret for the data we analyzed in this paper, and especially to Jens Koch for his patient guidance.

-
- [1] A. Blais. "Quantum information processing with circuit quantum electrodynamics." arXiv:cond-mat/0402216v1 (7 February 2004).
 - [2] R. H. Byrd, P. Lu, and J. Nocedal. "A Limited Memory Algorithm for Bound Constrained Optimization." *SIAM Journal on Scientific and Statistical Computing* 16, 5, pp. 1190-1208 (1995).
 - [3] R. P. Feynman, R. B. Leighton, and M. Sands. *The Feynman Lectures on Physics*. Addison-Wesley Publishing Company. Vol. 3, Ch. 21-9 (1963).
 - [4] Y. Nakamura et al. "Coherent control of macroscopic quantum states in a single-Cooper-pair box." *Nature* 398, pp. 786-788 (29 April 1999).
 - [5] L. K. Grover. "A fast quantum mechanical algorithm for database search." *Proceedings of the twenty-eighth annual ACM symposium on theory of computing*, pp. 212-219 (1996).
 - [6] J. Koch et al. "Charge insensitive qubit design derived from the Cooper pair box." arXiv:cond-mat/0703002v2 (28 February 2007).
 - [7] P. W. Shor. "Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer." *SIAM Review*, Vol. 41, No. 2, pp. 303-332 (June 1999).
 - [8] L. Vandersypen et al. "Experimental realization of Shor's quantum factoring algorithm using nuclear magnetic resonance." *Nature* 414, pp. 883-887 (20 December 2001).
 - [9] D. Vion et al. "Manipulating the Quantum State of an Electrical Circuit." *Science*, Vol. 3, No. 5569, pp. 886-889 (3 May 2002).
 - [10] A. Wallraff et al. "Strong coupling of a single photon to a superconducting qubit using circuit quantum electrodynamics." *Nature* 431, pp. 162-167 (9 September 2004).
 - [11] C. Zhu, R. H. Byrd, and J. Nocedal. "L-BFGS-B, FORTRAN routines for large scale bound constrained optimization." *ACM Transactions on Mathematical Software* Vol 23, Num. 4, pp. 550 - 560 (1997).